

Transactions

The Setting

- Database systems are normally being accessed by many users or processes at the same time. (Both queries and modifications.)
- Unlike operating systems, which support interaction of processes, a DBMS needs to keep processes from troublesome interactions.

Example: Bad Interaction

- You and your domestic partner each take \$100 from different ATM's at about the same time.
- The DBMS better make sure one account deduction doesn't get lost.
- Compare: An OS allows two people to edit a document at the same time. If both write, one's changes get lost.

ACID Transactions

- A DBMS is expected to support “ACID transactions,” processes that are:
 - – **Atomic** : Either the whole process is done or none is.
 - – **Consistent** : Database constraints are preserved.
 - – **Isolated** : It appears to the user as if only one process executes at a time.
 - – **Durable** : Effects of a process do not get lost if the system crashes.

Transactions in SQL

- SQL supports transactions, often behind the scenes.
 - – Each statement issued at the generic query interface is a transaction by itself.
 - – In programming interfaces like Embedded SQL or PSM, a transaction begins the first time a SQL statement is executed and ends with the program or an explicit transaction-end.

COMMIT

- The SQL statement COMMIT causes a transaction to complete.
- – It's database modifications are now permanent in the database.

ROLLBACK

- The SQL statement ROLLBACK also causes the transaction to end, but by **aborting**.
 - – No effects on the database.
- Failures like division by 0 or a constraint violation can also cause rollback, even if the programmer does not request it.

An Example: Interacting Processes

- Assume the usual **Sells(shop,product,price)** relation, and suppose that Joe's Shop sells only Coke for \$2.50 and Pepsi for \$3.00.
- Sally is querying **Sells** for the highest and lowest price Joe charges.
- Joe decides to stop selling Coke and Pepsi, but to sell only Fanta at \$3.50.

Sally's Program

- Sally executes the following two SQL statements, which we call **(min)** and **(max)**, to help remember what they do.
- **(max)** `SELECT MAX(price) FROM Sells WHERE shop = 'Joe's Shop';`
- **(min)** `SELECT MIN(price) FROM Sells WHERE shop = 'Joe's Shop';`

Joe's Program

- At about the same time, Joe executes the following steps, which have the mnemonic names **(del)** and **(ins)**.
- **(del)** DELETE FROM Sells WHERE shop = 'Joe's Shop';
- **(ins)** INSERT INTO Sells VALUES('Joe's Shop', 'Fanta', 3.50);

Interleaving of Statements

- Although (max) must come before (min), and (del) must come before (ins), there are no other constraints on the order of these statements, unless we group Sally's and/or Joe's statements into transactions.

Example: Strange Interleaving

- Suppose the steps execute in the order (max)(del)(ins)(min).
- Joe's Prices: 2.50, 3.00 2.50, 3.00 3.50
- Statement: (max) (del) (ins) (min)
- Result: 3.00 3.50
- Sally sees MAX < MIN!

Fixing the Problem by Using Transactions

- If we group Sally's statements (max)(min) into one transaction, then she cannot see this inconsistency.
- She sees Joe's prices at some fixed time.
 - – Either before or after he changes prices, or in the middle, but the MAX and MIN are computed from the same prices.

Another Problem: Rollback

- Suppose Joe executes `(del)(ins)`, not as a transaction, but after executing these statements, thinks better of it and issues a ROLLBACK statement.
- If Sally executes her statements after `(ins)` but before the rollback, she sees a value, 3.50, that never existed in the database.

Solution

- If Joe executes **(del)(ins)** as a transaction, its effect cannot be seen by others until the transaction executes COMMIT.
 - – If the transaction executes ROLLBACK instead, then its effects can never be seen.

Isolation Levels

- SQL defines four **isolation levels** = choices about what interactions are allowed by transactions that execute at about the same time.
- How a DBMS implements these isolation levels is highly complex, and a typical DBMS provides its own options.

Choosing the Isolation Level

- Within a transaction, we can say: SET TRANSACTION ISOLATION LEVEL X
- where X =
 - 1. SERIALIZABLE
 - 2. REPEATABLE READ
 - 3. READ COMMITTED
 - 4. READ UNCOMMITTED

Serializable Transactions

- If Sally = (max)(min) and Joe = (del)(ins) are each transactions, and Sally runs with isolation level SERIALIZABLE, then she will see the database either before or after Joe runs, but not in the middle.
- It's up to the DBMS vendor to figure out how to do that, e.g.:
 - – True isolation in time.
 - – Keep Joe's old prices around to answer Sally's queries.

Isolation Level Is Personal Choice

- Your choice, e.g., run serializable, affects only how you see the database, not how others see it.
- **Example:** If Joe Runs serializable, but Sally doesn't, then Sally might see no prices for Joe's Shop.
 - – i.e., it looks to Sally as if she ran in the middle of Joe's transaction.

Read-Committed Transactions

- If Sally runs with isolation level READ COMMITTED, then she can see only committed data, but not necessarily the same data each time.
- Example: Under READ COMMITTED, the interleaving (max)(del)(ins)(min) is allowed, as long as Joe commits.
 - – Sally sees $MAX < MIN$.

Repeatable-Read Transactions

- Requirement is like read-committed, plus: if data is read again, then everything seen the first time will be seen the second time.
 - – But the second and subsequent reads may see *more* tuples as well.

Example: Repeatable Read

- Suppose Sally runs under REPEATABLE READ, and the order of execution is (max) (del)(ins)(min).
 - – (max) sees prices 2.50 and 3.00.
 - – (min) can see 3.50, but must also see 2.50 and 3.00, because they were seen on the earlier read by (max).

Read Uncommitted

- A transaction running under READ UNCOMMITTED can see data in the database, even if it was written by a transaction that has not committed (and may never).
- Example: If Sally runs under READ UNCOMMITTED, she could see a price 3.50 even if Joe later aborts.

Syntax

```
BEGIN;  
UPDATE accounts SET balance = balance - 100.00  
|   WHERE name = 'Alice';  
-- etc etc  
COMMIT;
```


BEGIN

```
BEGIN [ WORK | TRANSACTION ] [ transaction_mode [ , ... ] ]
```

where *transaction_mode* is one of:

```
ISOLATION LEVEL { SERIALIZABLE | REPEATABLE READ | READ  
COMMITTED | READ UNCOMMITTED }  
READ WRITE | READ ONLY
```

Savepoint

```
BEGIN;  
UPDATE accounts SET balance = balance - 100.00  
|   WHERE name = 'Alice';  
SAVEPOINT my_savepoint;  
UPDATE accounts SET balance = balance + 100.00  
|   WHERE name = 'Bob';  
-- oops ... forget that and use Wally's account  
ROLLBACK TO my_savepoint;  
UPDATE accounts SET balance = balance + 100.00  
|   WHERE name = 'Wally';  
COMMIT;
```

Savepoint

```
BEGIN;  
    INSERT INTO table1 VALUES (1);  
    SAVEPOINT my_savepoint;  
    INSERT INTO table1 VALUES (2);  
    ROLLBACK TO SAVEPOINT my_savepoint;  
    INSERT INTO table1 VALUES (3);  
COMMIT;
```

Rollback

```
BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;  
UPDATE accounts SET balance = balance - 100.00  
|   WHERE name = 'Alice';  
SAVEPOINT my_savepoint;  
UPDATE accounts SET balance = balance + 100.00  
|   WHERE name = 'Bob';  
-- oops ... forget that and use Wally's account  
ROLLBACK TO my_savepoint;  
UPDATE accounts SET balance = balance + 100.00  
|   WHERE name = 'Wally';  
COMMIT;
```

Release Savepoint

```
BEGIN;  
    INSERT INTO table1 VALUES (3);  
    SAVEPOINT my_savepoint;  
    INSERT INTO table1 VALUES (4);  
    RELEASE SAVEPOINT my_savepoint;  
COMMIT;
```

Questions